

MaxMiner: Mining Max-patterns

- Review!
 - An itemset X is a **max-pattern** if X is frequent and there exists no frequent super-pattern $Y \supset X$
 - i.e., no such a Y
 - Y is a super-pattern of X
 - The support of Y is greater than minSup
 - The support of Y can be smaller than that of X
- MaxMiner is based on the Apriori algorithm
- R. Bayardo. Efficiently mining long patterns from databases. In *SIGMOD'98*



MaxMiner: Mining Max-patterns

- 1st scan: find frequent items and **sort** them (**ascending order**)
 - A, B, C, D, E (E is most frequently occurring)
- 2nd scan: find support for 2-itemsets with max-patterns
 - **AB**, AC, AD, AE, **ABCDE** ←
 - **BC**, BD, BE, **BCDE** ←
 - **CD**, CE, **CDE** ←
 - **DE**

Potential max-patterns

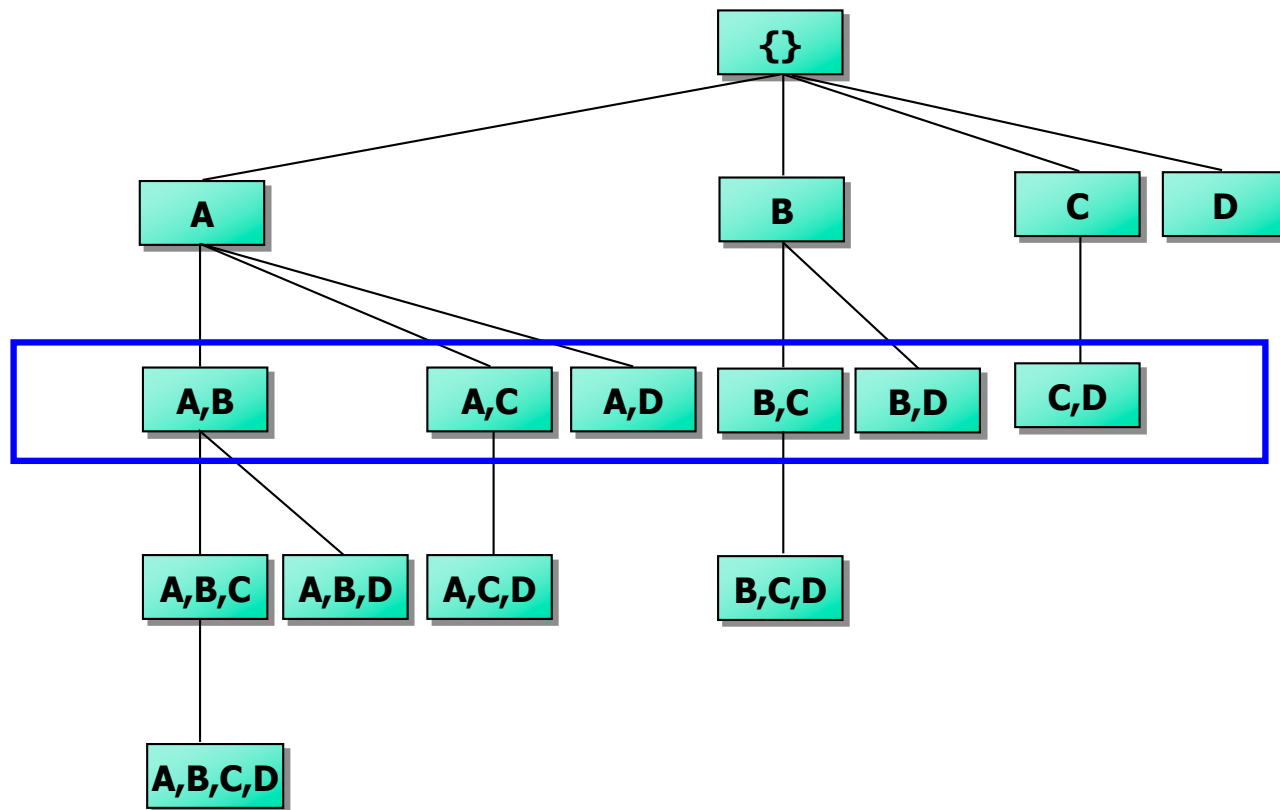
Tid	Items
10	A,B,C,D,E
20	B,C,D,E,
30	A,C,D,F,E

- Reduce a lot of candidates in later stages
 - Since BCDE is a max-pattern, no need to check BCD, BDE, CDE in later scan
 - If AC is infrequent, no need to check ABC in later scans



MaxMiner: Mining Max-patterns

Complete *set-enumeration tree* over four items



MaxMiner: Mining Max-patterns

- 1st scan: find frequent items and **sort** them (**ascending order**)
 - A, B, C, D, E (E is most frequently occurring)
- 2nd scan: find support for 2-itemsets with max-patterns
 - **AB**, AC, AD, AE, **ABCDE** ←
 - **BC**, BD, BE, **BCDE** ←
 - **CD**, CE, **CDE** ←
 - **DE**

Potential max-patterns

Tid	Items
10	A,B,C,D,E
20	B,C,D,E,
30	A,C,D,F,E

- Reduce a lot of candidates in later stages
 - Since BCDE is a max-pattern, no need to check BCD, BDE, CDE in later scan
 - If AC is infrequent, no need to check ABC in later scans



Mining Closed Patterns: CLOSET

- Review
 - An itemset X is **closed** if X is *frequent* and there exists *no super-pattern* $Y \supset X$, *with the same support as X*
 - i.e., no such a Y
 - Y is a super-pattern of X
 - The support of Y is should be the same as that of X



Mining Closed Patterns: CLOSET

- Use the FP-tree for finding frequent patterns
- Flist: list of all frequent items in a support descending order
 - Flist: c-e-f-a-d
- Divide search space
 - Patterns having d
 - Patterns having a but no d
 - Patterns having f but no d and a
 - Patterns having e but no d, a, and f
 - Patterns having c but no d, a, f, and e

Min_sup=2

TID	Items
10	a, c, d, e, f
20	a, b, e
30	c, e, f
40	a, c, d, f
50	c, e, f



Mining Closed Patterns: CLOSET

- Naïve approach: Quite costly!
 - To mine a complete set of all frequent itemsets
 - To remove every frequent itemset whose support is the same as that of its superset
- Find only the closed itemsets recursively in an efficient way during the mining process using the FP-tree
 - Key idea: every transaction having d also has $cfa \Rightarrow cfad$ is a frequent closed pattern
 - You can consider details by referring to FP-Growth
- J. Pei, J. Han & R. Mao. "CLOSET: An Efficient Algorithm for Mining Frequent Closed Itemsets", DMKD'00



CHARM: Mining by Exploring Vertical Data Format

- Vertical format: $t(AB) = \{T_{11}, T_{25}, \dots\}$
 - tid-list: list of trans.-ids containing an itemset
- Algorithm
 - Transform a horizontally formatted data to a vertically format by scanning the dataset once
 - Easy: # of items is much smaller than that of transactions
 - Starting with $k=1$, construct candidate $(k+1)$ -itemsets from frequent k -itemsets
 - Using the TID-sets intersection and Apriori property
 - Repeat this process with k incremented by 1 until no frequent itemsets can be found



CHARM: Mining by Exploring Vertical Data Format

Table 5.3 The vertical data format of the transaction data set D of Table 5.1.

<i>itemset</i>	<i>TID_set</i>
I1	{T100, T400, T500, T700, T800, T900}
I2	{T100, T200, T300, T400, T600, T800, T900}
I3	{T300, T500, T600, T700, T800, T900}
I4	{T200, T400}
I5	{T100, T800}



CHARM: Mining by Exploring Vertical Data Format

Table 5.4 The 2-itemsets in vertical data format.

itemset	TID_set
{I1, I2}	{T100, T400, T800, T900}
{I1, I3}	{T500, T700, T800, T900}
{I1, I4}	{T400}
{I1, I5}	{T100, T800}
{I2, I3}	{T300, T600, T800, T900}
{I2, I4}	{T200, T400}
{I2, I5}	{T100, T800}
{I3, I5}	{T800}

Table 5.5 The 3-itemsets in vertical data format.

itemset	TID_set
{I1, I2, I3}	{T800, T900}
{I1, I2, I5}	{T100, T800}

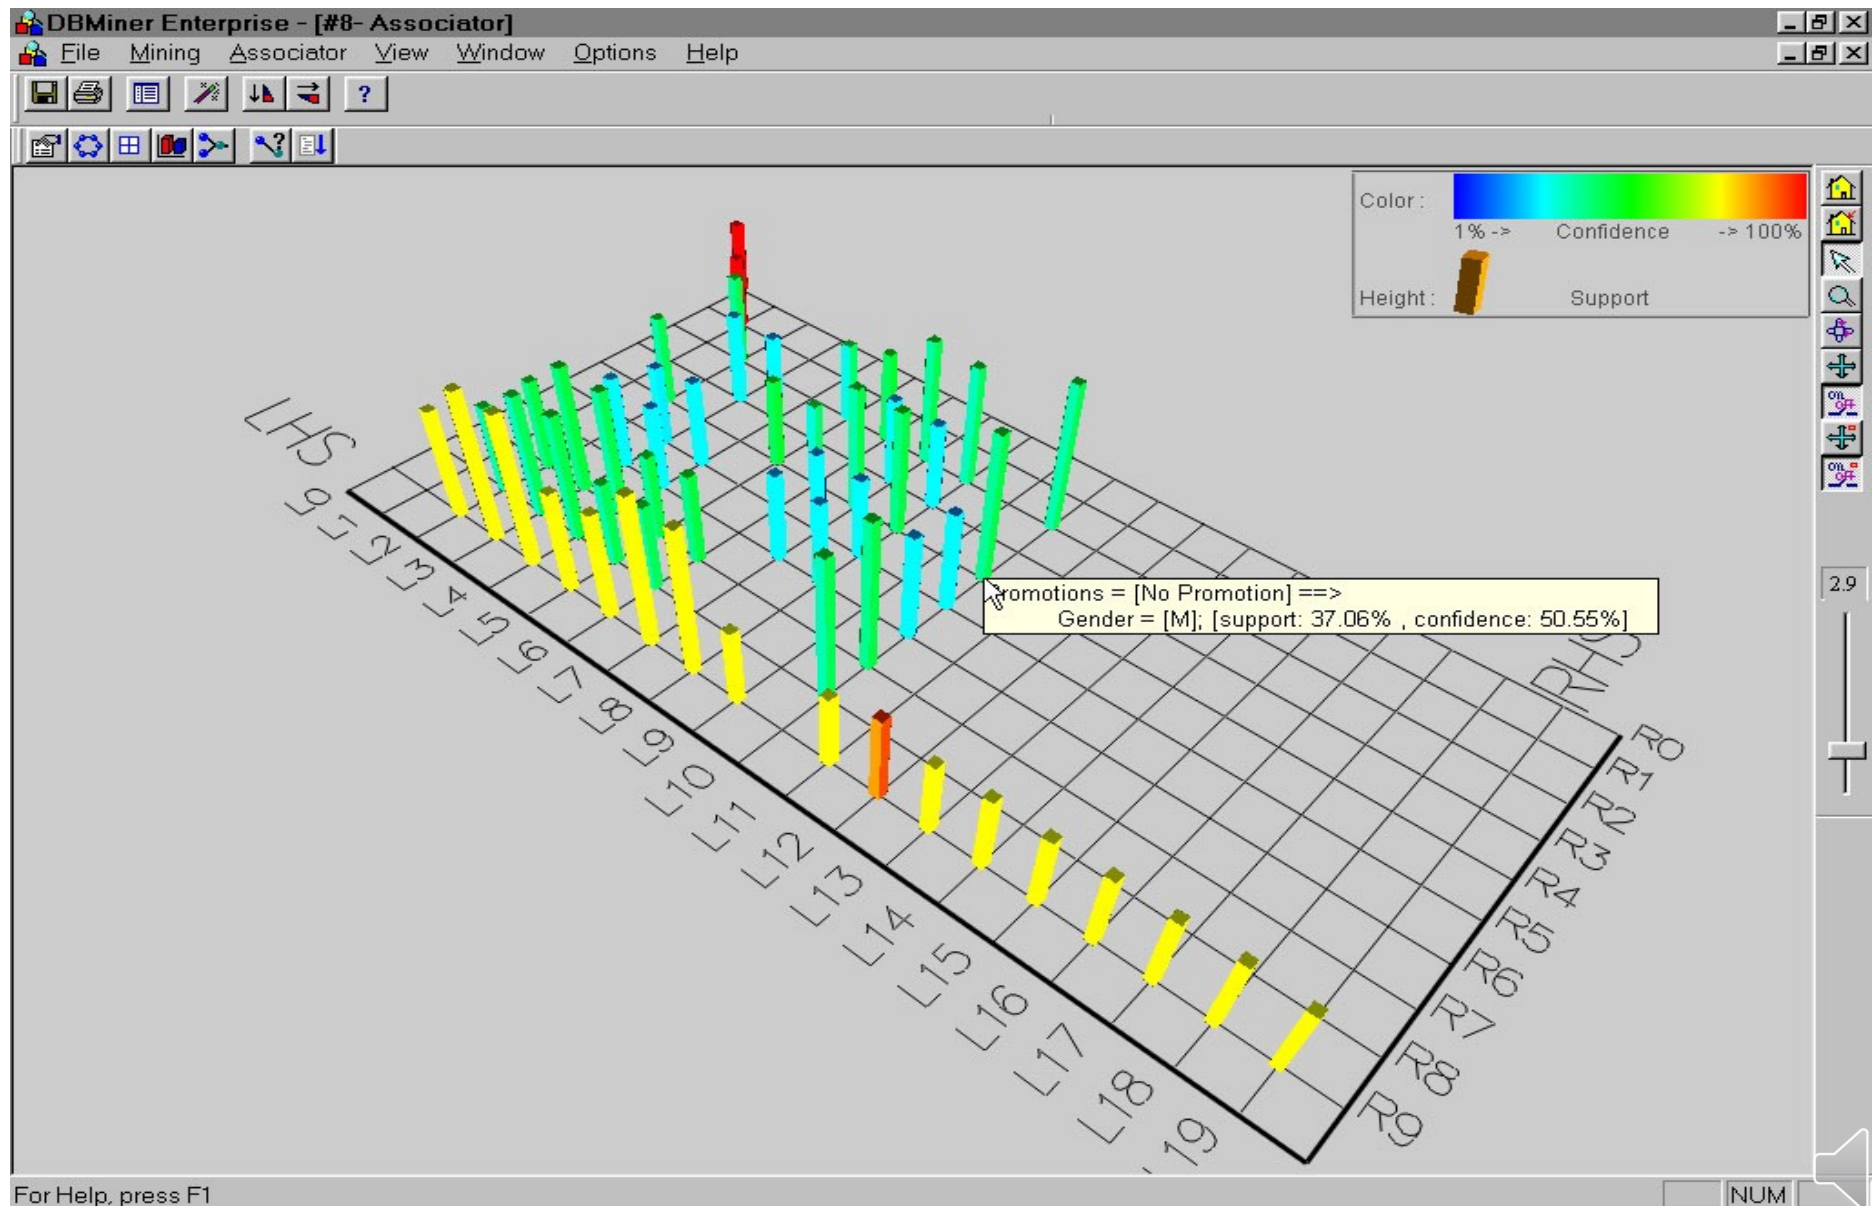


CHARM: Mining by Exploring Vertical Data Format

- No need to scan a database to find the support of $(k+1)$ itemsets
 - TID set of each k -itemset carries sufficient information including a support value
 - But, it is quite long and requires large space for intersection



Visualization of Association Rules: Plane Graph



Chapter 5: Mining Frequent Patterns, Association and Correlations

- Basic concepts and a road map
- Efficient and scalable frequent itemset mining methods
- Mining various kinds of association rules
- From association mining to correlation analysis
- Constraint-based association mining
- Summary



Mining Various Kinds of Association Rules

- Mining multilevel association
- Mining multidimensional association
- Mining quantitative association
- Mining interesting correlation patterns



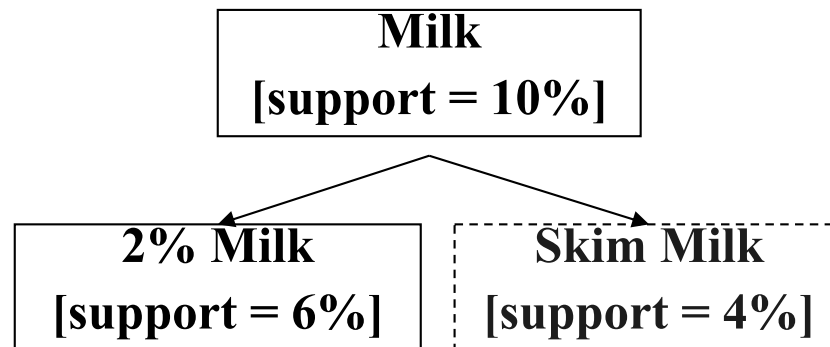
Mining Multiple-Level Association Rules

- Items often form hierarchies
- Flexible support settings
 - Items at the lower level are expected to have lower support
- Exploration of *shared* multi-level mining (Agrawal & Srikant@VLDB'95, Han & Fu@VLDB'95)

uniform support

Level 1
min_sup = 5%

Level 2
min_sup = 5%



reduced support

Level 1
min_sup = 5%

Level 2
min_sup = 3%



Multi-level Association: Redundancy Filtering

- Some rules may be redundant due to “ancestor” relationships between items.
- Example
 - milk \Rightarrow wheat bread [support = 8%, confidence = 70%]
 - 2% milk \Rightarrow wheat bread [support = 2%, confidence = 72%]
- We say the first rule is an ancestor of the second rule.
- A (descendent) rule is redundant if
 - Its **support** is close to the *“expected” value*, based on the rule’s ancestor
 - Its **confidence** is close to that of the rule’s ancestor



Mining Multi-Dimensional Association

- Single-dimensional rules: (having a dimension or a predicate)

$\text{buys}(X, \text{"milk"}) \Rightarrow \text{buys}(X, \text{"bread"})$: milk \Rightarrow bread

- Multi-dimensional rules: ≥ 2 dimensions or predicates

- Inter-dimension assoc. rules (*no repeated predicates*)

$\text{age}(X, \text{"19-25"}) \wedge \text{occupation}(X, \text{"student"}) \Rightarrow \text{buys}(X, \text{"coke"})$

- hybrid-dimension assoc. rules (*repeated predicates*)

$\text{age}(X, \text{"19-25"}) \wedge \text{buys}(X, \text{"popcorn"}) \Rightarrow \text{buys}(X, \text{"coke"})$



Attribute Types

- $\text{age}(X, "19-25") \wedge \text{occupation}(X, "student") \Rightarrow \text{buys}(X, "coke")$
attributes
- Categorical Attributes
 - Finite number of possible values, no ordering among values
- Quantitative Attributes
 - Numeric, implicit ordering among values
 - Discretization and clustering approaches required (**why?**)



Mining Quantitative Associations

- Techniques can be categorized by how numerical attributes, such as **age** or **salary** are treated
 1. Static discretization based on predefined concept hierarchies (data cube methods)
 2. Dynamic discretization based on data distribution (quantitative rules, e.g., Agrawal & Srikant@SIGMOD96)
 3. Clustering: Distance-based association (e.g., Yang & Miller@SIGMOD97)



Static Discretization of Quantitative Attributes

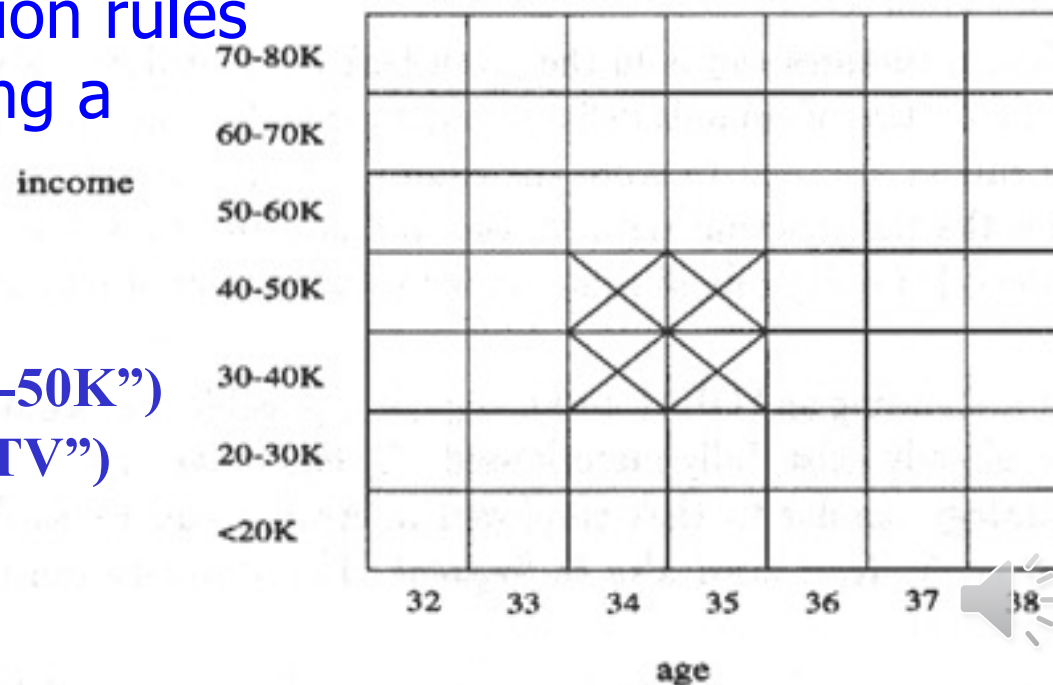
- Discretized prior to mining using a concept hierarchy
 - $\text{age}(X, "19-25") \wedge \text{occupation}(X, "student") \Rightarrow \text{buys}(X, "coke")$
 - Numeric values are replaced by ranges (as a categorical value)
- In a relational database, finding all frequent k -predicate sets will require k or $k+1$ table scans.

Quantitative Association Rules

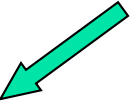
- Proposed by Lent, Swami and Widom ICDE'97
- Numeric attributes are *dynamically* discretized
- 2-D quantitative association rules: $A_{\text{quan1}} \wedge A_{\text{quan2}} \Rightarrow A_{\text{cat}}$
 - The confidence is higher than threshold
 - The support is higher than threshold
- Cluster *adjacent* association rules to form general rules using a 2-D grid
- Example

$\text{age}(X, "34-35") \wedge \text{income}(X, "30-50K")$
 $\Rightarrow \text{buys}(X, "high\ resolution\ TV")$

Note: simplified!



Chapter 5: Mining Frequent Patterns, Association and Correlations

- Basic concepts and a road map
- Efficient and scalable frequent itemset mining methods
- Mining various kinds of association rules
- From association mining to correlation analysis 
- Constraint-based association mining
- Summary



Interestingness Measure: Correlations (Lift)

- *play basketball* \Rightarrow *eat cereal* [40%, 66.7%] is misleading
 - The overall % of students eating cereal is 75% > 66.7%.
- *play basketball* \Rightarrow *not eat cereal* [20%, 33.3%] is more meaningful, although it has lower support and confidence
- Measure of dependent/correlated events: **lift**

Contingency table

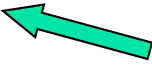
$$lift = \frac{P(A \cup B)}{P(A)P(B)}$$

	Basketball	Not basketball	Sum (row)
Cereal	2000	1750	3750
Not cereal	1000	250	1250
Sum(col.)	3000	2000	5000

$$lift(B, C) = \frac{2000 / 5000}{3000 / 5000 * 3750 / 5000} = 0.89 \quad lift(B, \neg C) = \frac{1000 / 5000}{3000 / 5000 * 1250 / 5000} = 1.33$$



Chapter 5: Mining Frequent Patterns, Association and Correlations

- Basic concepts and a road map
- Efficient and scalable frequent itemset mining methods
- Mining various kinds of association rules
- From association mining to correlation analysis
- Constraint-based association mining 
- Summary



Constraint-based (Query-Directed) Mining

- Finding **all** the patterns in a database **autonomously**? — unrealistic!
 - The patterns could be too many but not focused!
- Data mining should be an **interactive** process
 - User directs what to be mined
 - She/he uses a **data mining query language** (or a graphical user interface)
- Constraint-based mining
 - User flexibility: provides **constraints** on what to be mined
 - System optimization: explores such constraints for efficient mining—**constraint-based mining**



Constraints in Data Mining

- Knowledge type constraint:
 - classification, association, clustering etc.
- Data constraint — using SQL-like queries
 - find product pairs sold together in stores in **Chicago** in **Dec.'21**
- Dimension/level constraint
 - in relevance to **region, price, brand, customer category**
- Interestingness constraint
 - strong rules: $\text{min_support} \geq 3\%$, $\text{min_confidence} \geq 60\%$

Constrained Mining vs. Other Operations

- Constrained mining vs. constraint-based search
 - Both are aimed at reducing search space
 - Constrained mining: finding all patterns satisfying constraints
 - Constraint-based search: finding some (or one) answer in constraint-based search in AI



Constrained Mining vs. Other Operations

- Constrained mining vs. query processing in DBMS
 - Both are aimed at finding **all answers**
 - Query processing: finding **tuples** in a database
 - Constrained mining: discovering **patterns hidden** in a database
 - Constrained mining shares a similar philosophy as **pushing selections deeply** in query processing

Anti-Monotonicity in Constraint Pushing

- Anti-monotonicity on a constraint
 - When an itemset S **violates** the constraint, so does any of its superset
 - $\text{sum}(S.\text{Price}) \leq v$ is **anti-monotone**
 - $\text{sum}(S.\text{Price}) \geq v$ is **not anti-monotone**
- Example. C: $\text{range}(S.\text{profit}) \leq 15$ is **anti-monotone**
 - Itemset ab violates C
 - So does every superset of ab

TDB (min_sup=2)

TID	Transaction
10	a, b, c, d, f
20	b, c, d, f, g, h
30	a, c, d, e, f
40	c, e, f, g

Item	Profit
a	40
b	0
c	-20
d	10
e	-30
f	30
g	20
h	-10

Monotonicity for Constraint Pushing

TDB (min_sup=2)

- Monotonicity on a constraint
 - When an itemset S **satisfies** the constraint, so does any of its superset
 - $\text{sum}(S.\text{Price}) \geq v$ is **monotone**
 - $\text{min}(S.\text{Price}) \leq v$ is **monotone**
- Example. C: $\text{range}(S.\text{profit}) \geq 15$
 - Itemset ab satisfies C
 - So does every superset of ab

TID	Transaction
10	a, b, c, d, f
20	b, c, d, f, g, h
30	a, c, d, e, f
40	c, e, f, g

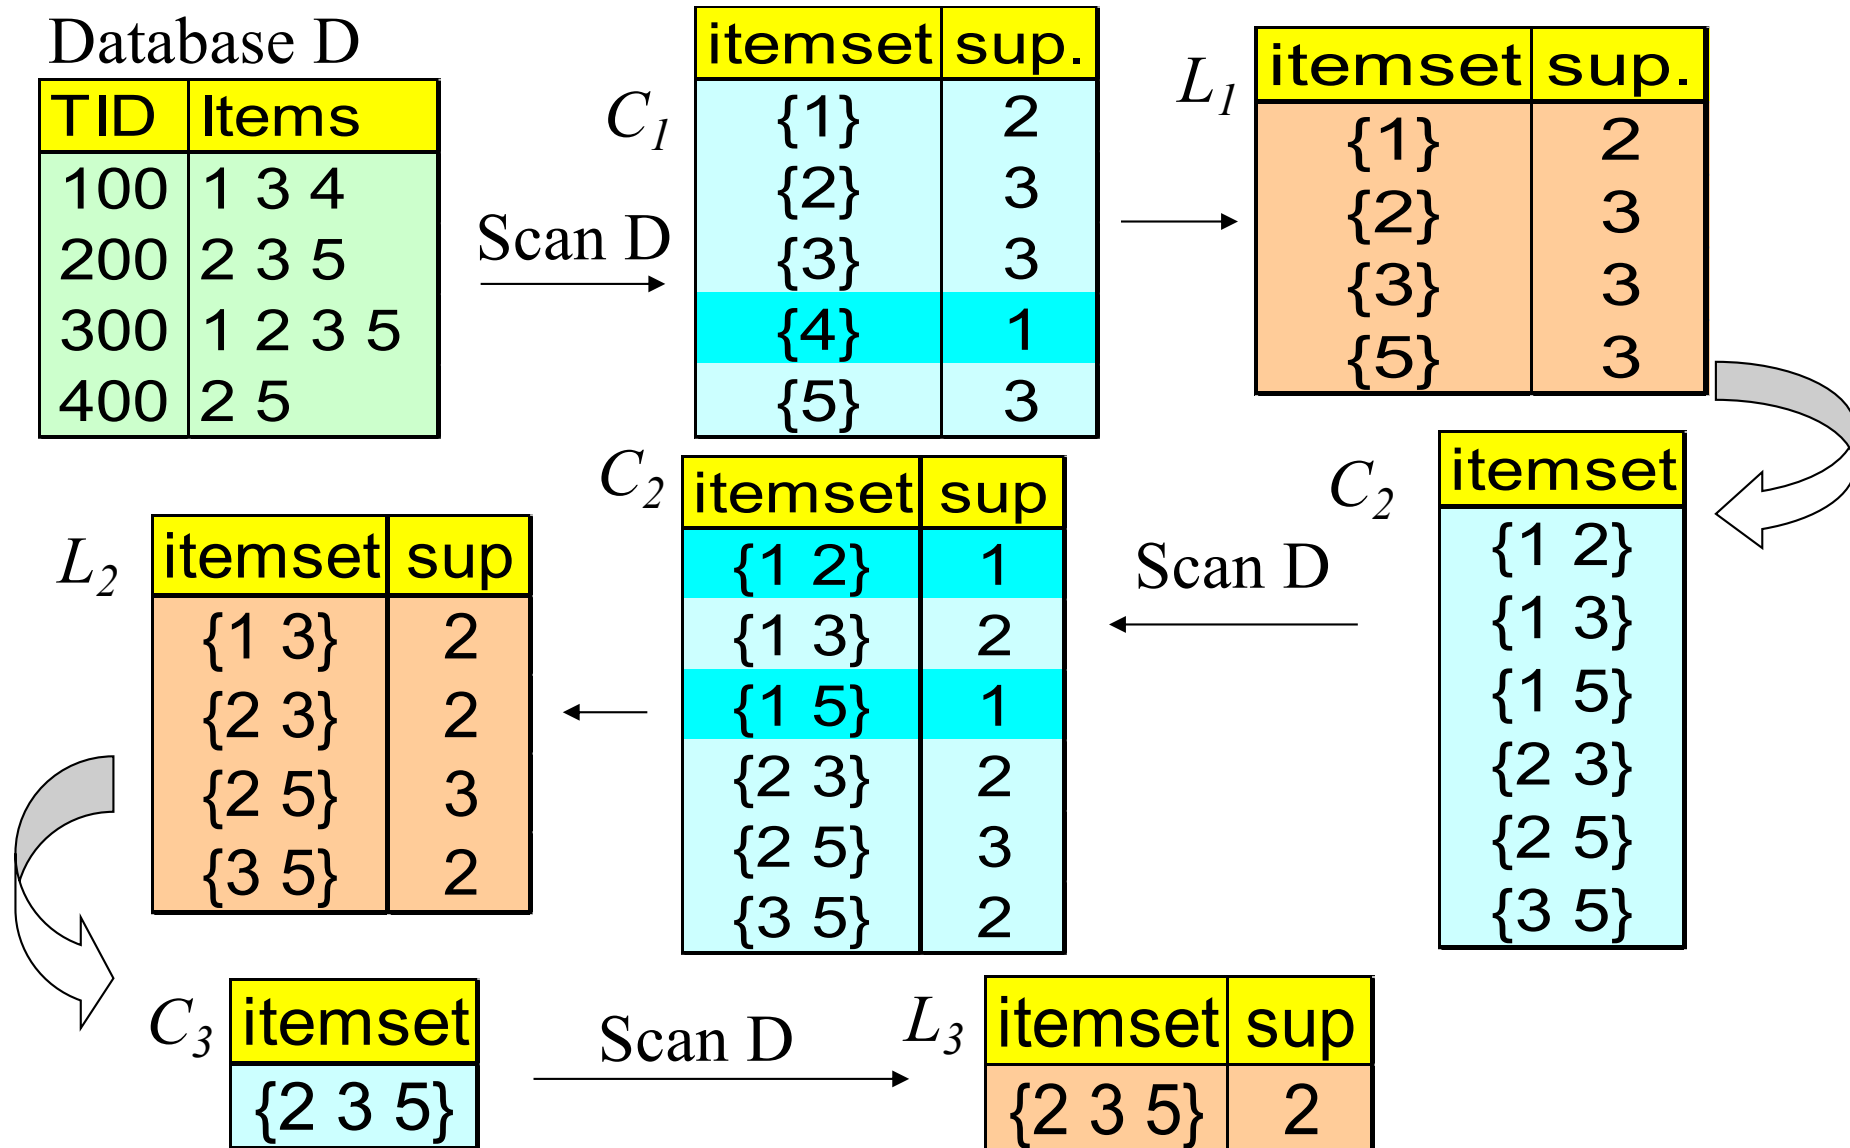
Item	Profit
a	40
b	0
c	-20
d	10
e	-30
f	30
g	20
h	-10

Succinctness

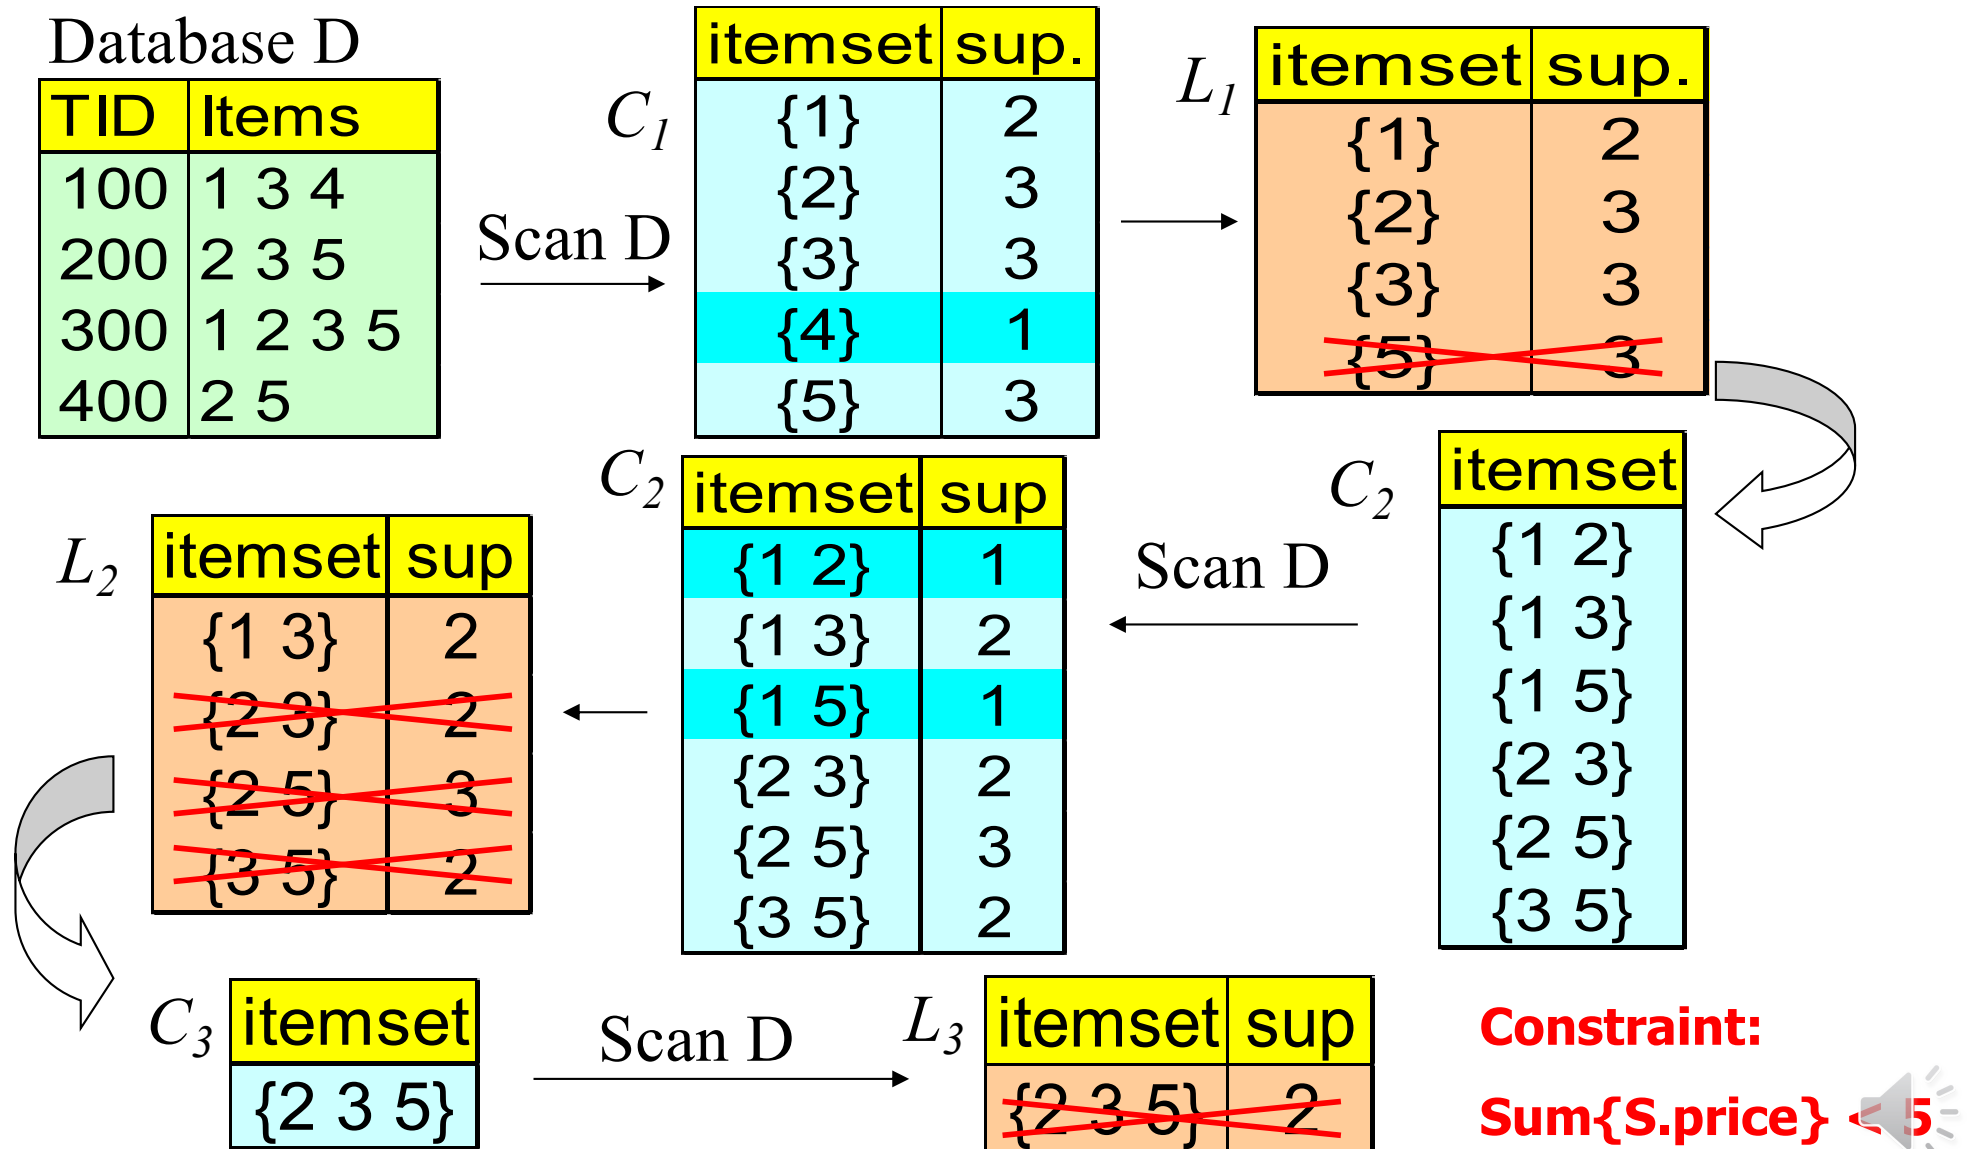
- Succinctness on a constraint
 - When a set of items (say A_I) satisfies a constraint C , any set S satisfying C is 'simply computed' based on A_I (In this case, S contains a subset belonging to A_I)
 - $\min(S.Price) \leq v$ is succinct
 - $\sum(S.Price) \geq v$ is not succinct
- Good thing
 - Without looking at the transaction database, whether an itemset S satisfies constraint C can be determined based on the selection of items
- Optimization: If C is succinct, C is pre-counting pushable



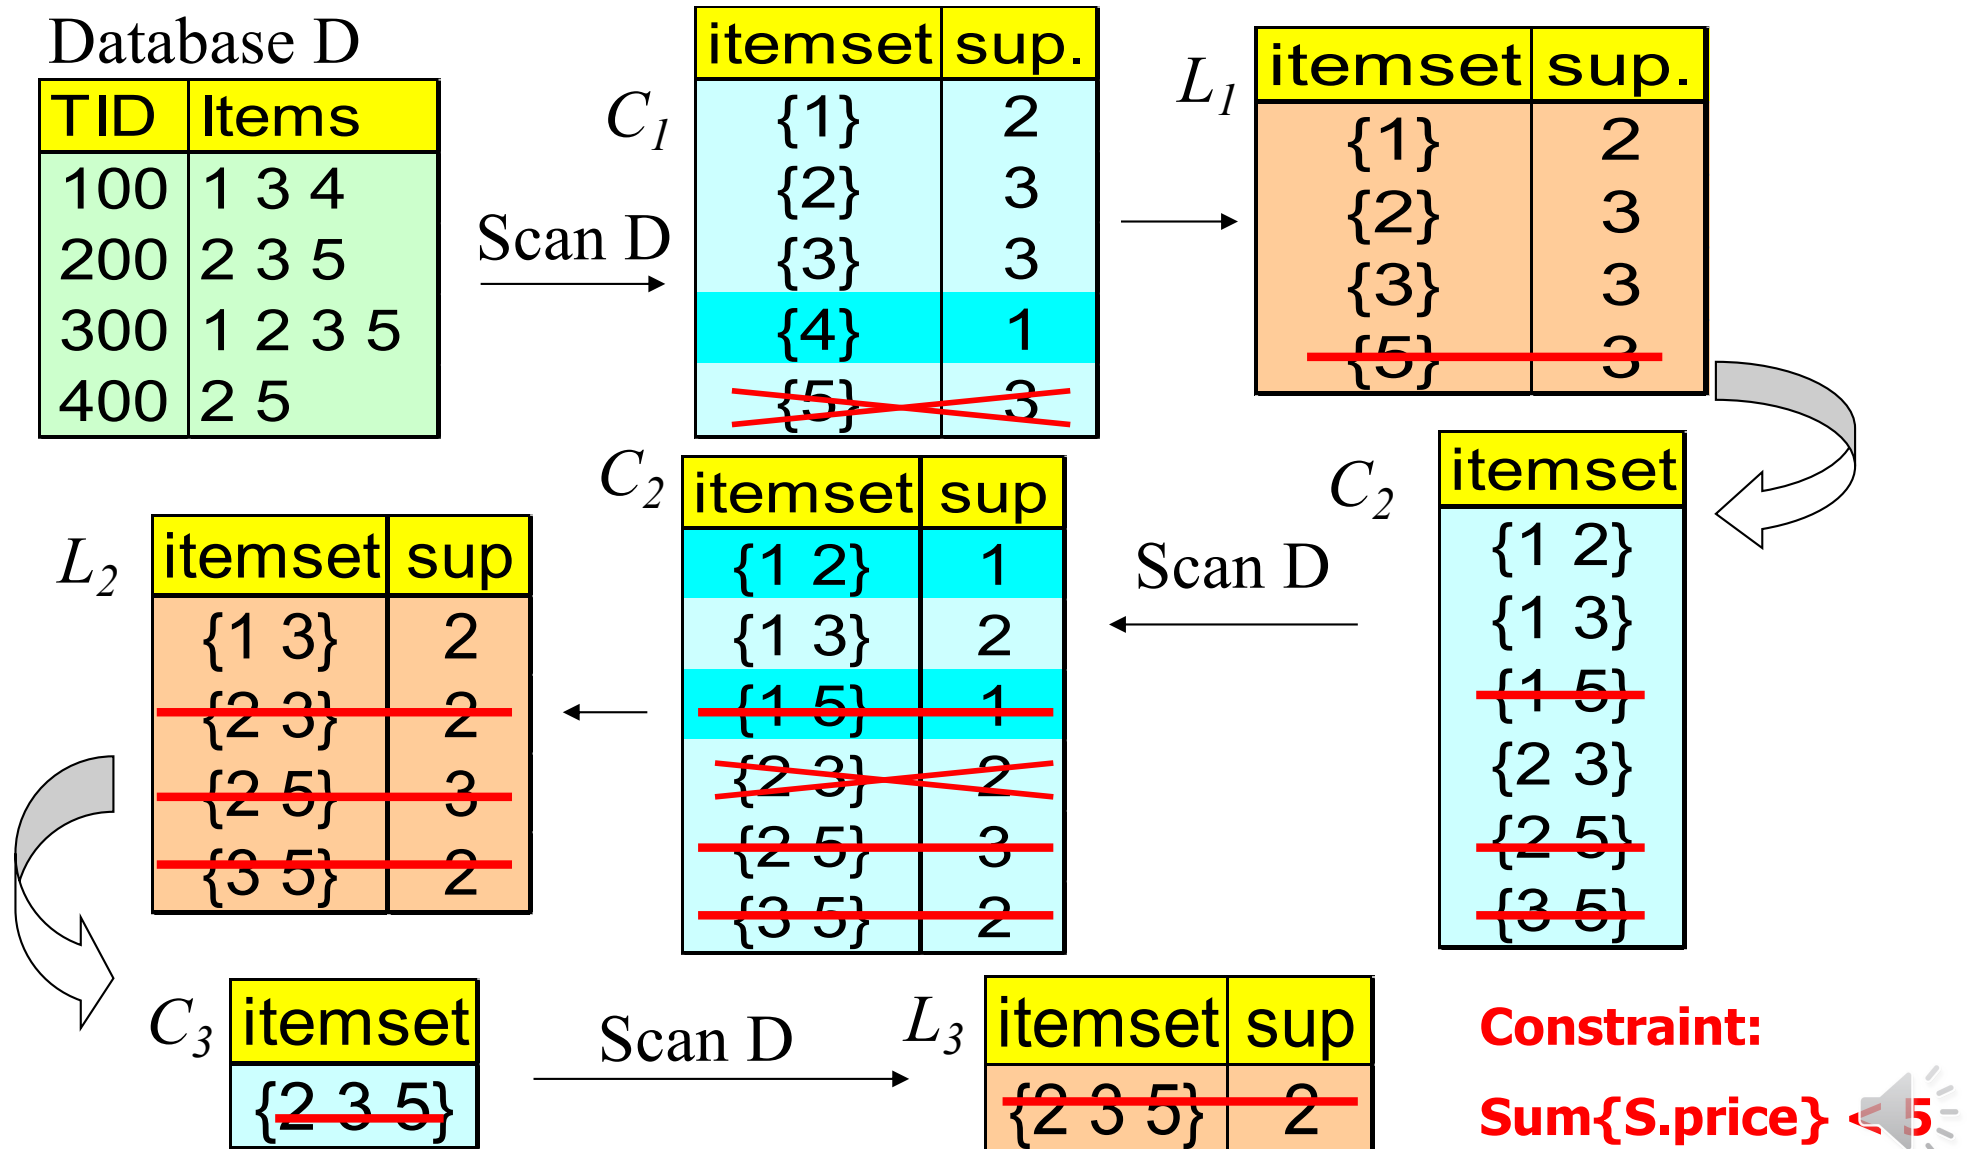
The Apriori Algorithm — Example



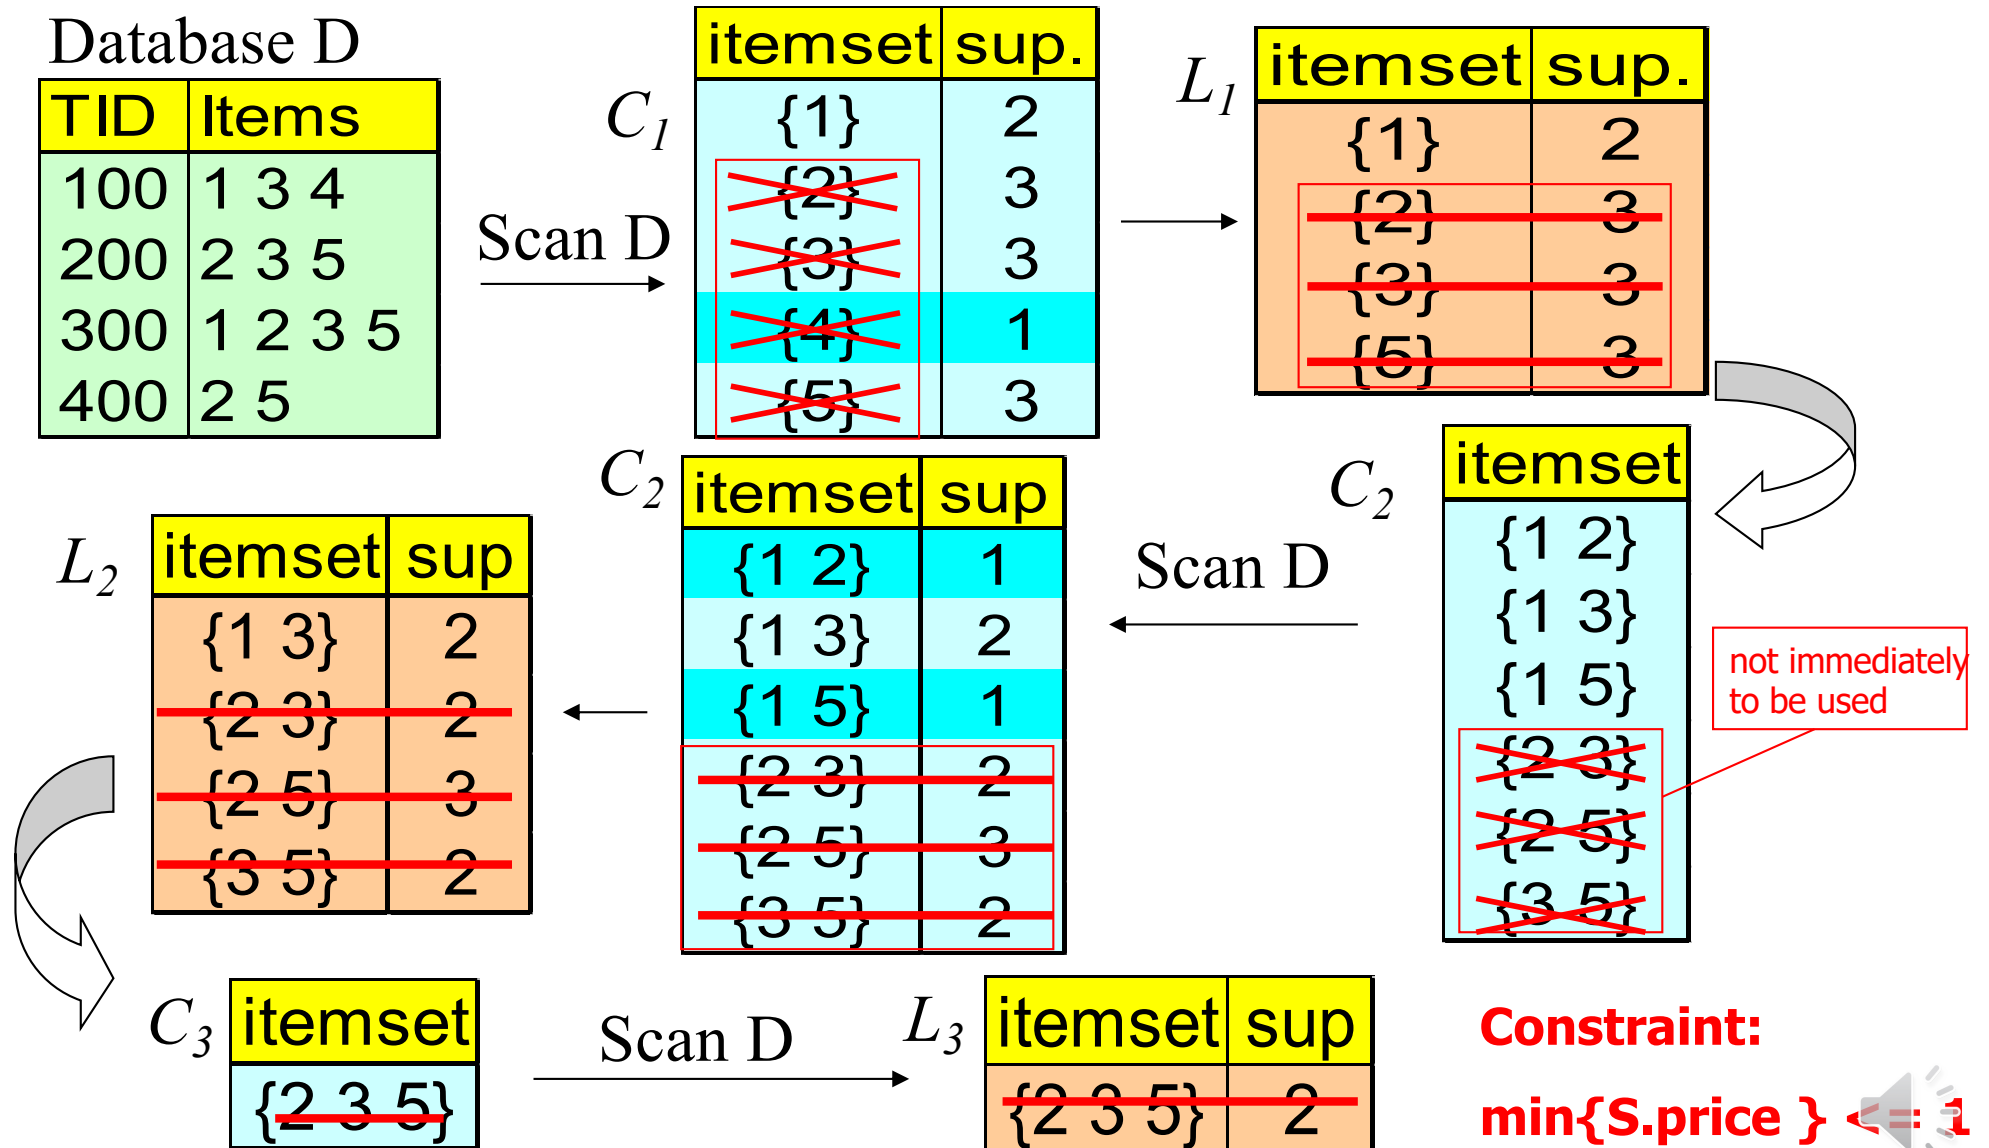
Naïve Algorithm: Apriori + Constraint



The Constrained Apriori Algorithm: Push an Anti-monotone Constraint Deep



The Constrained Apriori Algorithm: Push a Succinct Constraint Deep



Converting “Tough” Constraints

- Convert tough constraints into anti-monotone or monotone by properly ordering items
- Examine C: $\text{avg}(S.\text{profit}) \geq 25$
 - Order items in **value-descending** order
 - $\langle a, f, g, d, b, h, c, e \rangle$
 - If an itemset afb violates C
 - So does $afbh$
 - It becomes **anti-monotone!**

TDB (min_sup=2)

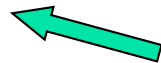
TID	Transaction
10	a, b, c, d, f
20	b, c, d, f, g, h
30	a, c, d, e, f
40	c, e, f, g

Item	Profit
a	40
b	0
c	-20
d	10
e	-30
f	30
g	20
h	-10



Chapter 5: Mining Frequent Patterns, Association and Correlations

- Basic concepts and a road map
- Efficient and scalable frequent itemset mining methods
- Mining various kinds of association rules
- From association mining to correlation analysis
- Constraint-based association mining
- Summary



Frequent-Pattern Mining: Summary

- Frequent pattern mining—an important task in data mining
- Scalable frequent pattern mining methods
 - Apriori (Candidate generation & test)
 - Projection-based (FPgrowth, CLOSET+, ...)
 - Vertical format approach (CHARM, ...)
- Mining a variety of rules and interesting patterns
- Constraint-based mining
- Extensions and applications



