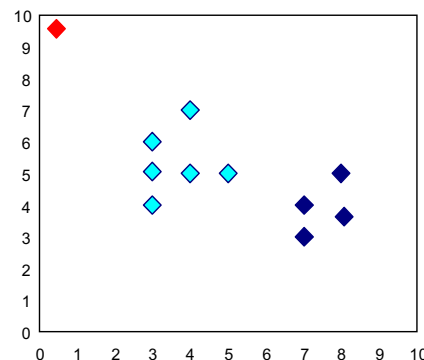


What Is the Problem of the K-Means Method?

- The k-means algorithm is sensitive to *outliers* !
 - An object with an extremely large value may substantially distort the distribution of the data
- K-Medoids: Instead of taking the **mean** value (i.e., *centroids*) of the object in a cluster as a reference point, *a medoids* can be used, which is the *most centrally-located object* in a cluster



The *K-Medoids* Clustering Method

- Find *representative* objects, called medoids, in clusters
 - *PAM* (Partitioning Around Medoids, 1987)
 - *CLARA* (Kaufmann & Rousseeuw, 1990)
 - *CLARANS* (Ng & Han, 1994): Randomized sampling

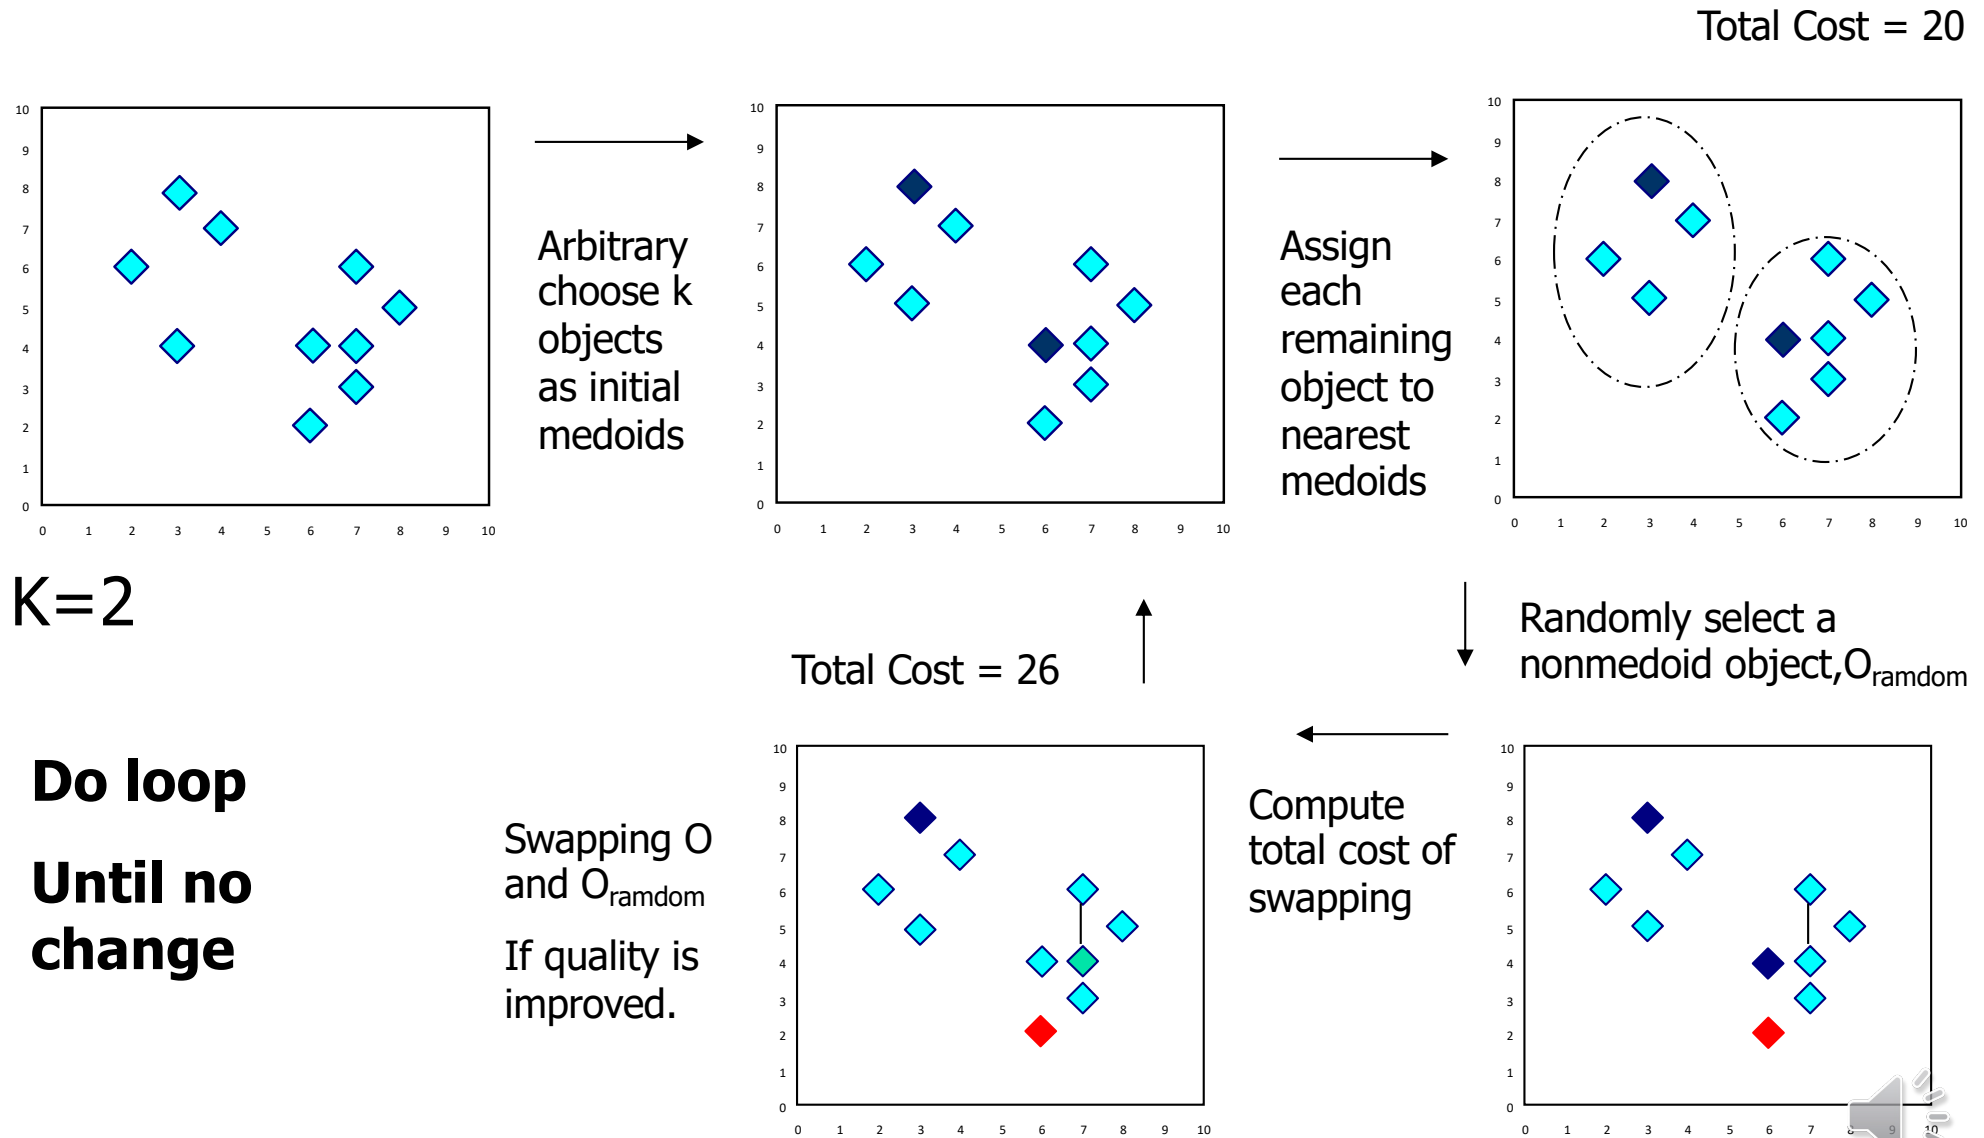


PAM (Partitioning Around Medoids) (1987)

- PAM (Kaufman and Rousseeuw, 1987), built in Splus
- Use a **real object** to represent the cluster
 - Select k representative objects arbitrarily
 - For each pair of non-selected object h and **selected object (i.e., seed) i** , calculate the total swapping cost TC_{ih}
 - For each pair of i and h ,
 - If $TC_{ih} < 0$, i is replaced by h
 - Then, each non-selected object is assigned to the most similar representative object
 - Repeat steps 2-3 until there is no change



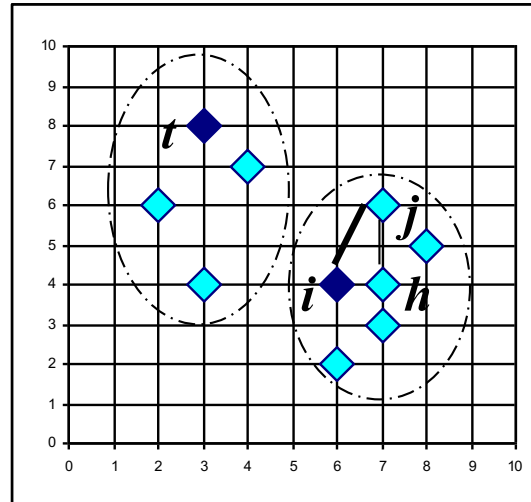
A Typical K-Medoids Algorithm (PAM)



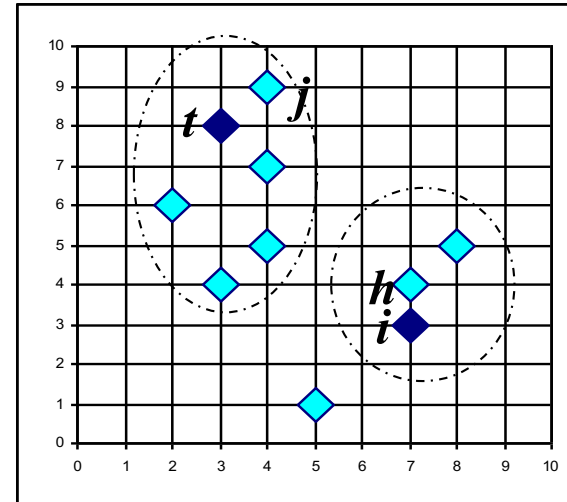
PAM Clustering: Total swapping cost $TC_{ih} = \sum_j C_{jih}$

NewC - OldC

i: original seed
h: new seed
t: other seed
j: non-seed

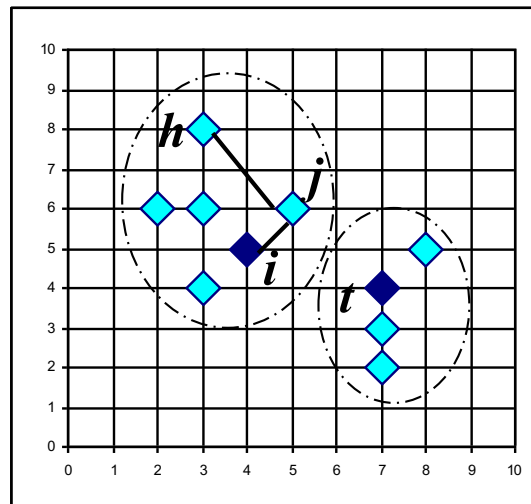


$$C_{jih} = d(j, h) - d(j, i)$$

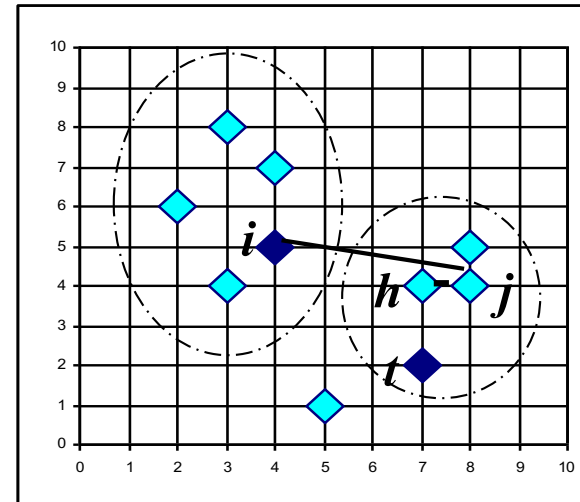


$$C_{jih} = 0$$

A: j belonged to i and now belongs to h
B: j belonged to t and again belongs to t
C: j belonged to i and now belongs to t
D: j belonged to t and now belongs to h



$$C_{jih} = d(j, t) - d(j, i)$$



$$C_{jih} = d(j, h) - d(j, t)$$



What Is the Problem with PAM?

- PAM is more robust than k-means in the presence of noise and outliers
 - because a medoid is less influenced by outliers or other extreme values than a mean (i.e., centroid)
- PAM works efficiently for small data sets but **does not scale well** for large data sets.
 - $O(i * k * (n - k)^2)$ where n is # of data, k is # of clusters, i is # of iterations

→ Sampling based method,

CLARA (Clustering LARge Applications)



CLARA (Clustering Large Applications) (1990)

- CLARA (Kaufmann and Rousseeuw in 1990)
 - Built in statistical analysis packages, such as S+
- It draws *multiple samples* of the data set, applies *PAM* on each sample, and gives the best clustering as the output
- Strength: deals with larger data sets than *PAM*
- Weakness:
 - Efficiency *depends on the sample size*
 - A good clustering based on samples will not necessarily represent a good clustering of the whole data set *if the sample is biased*



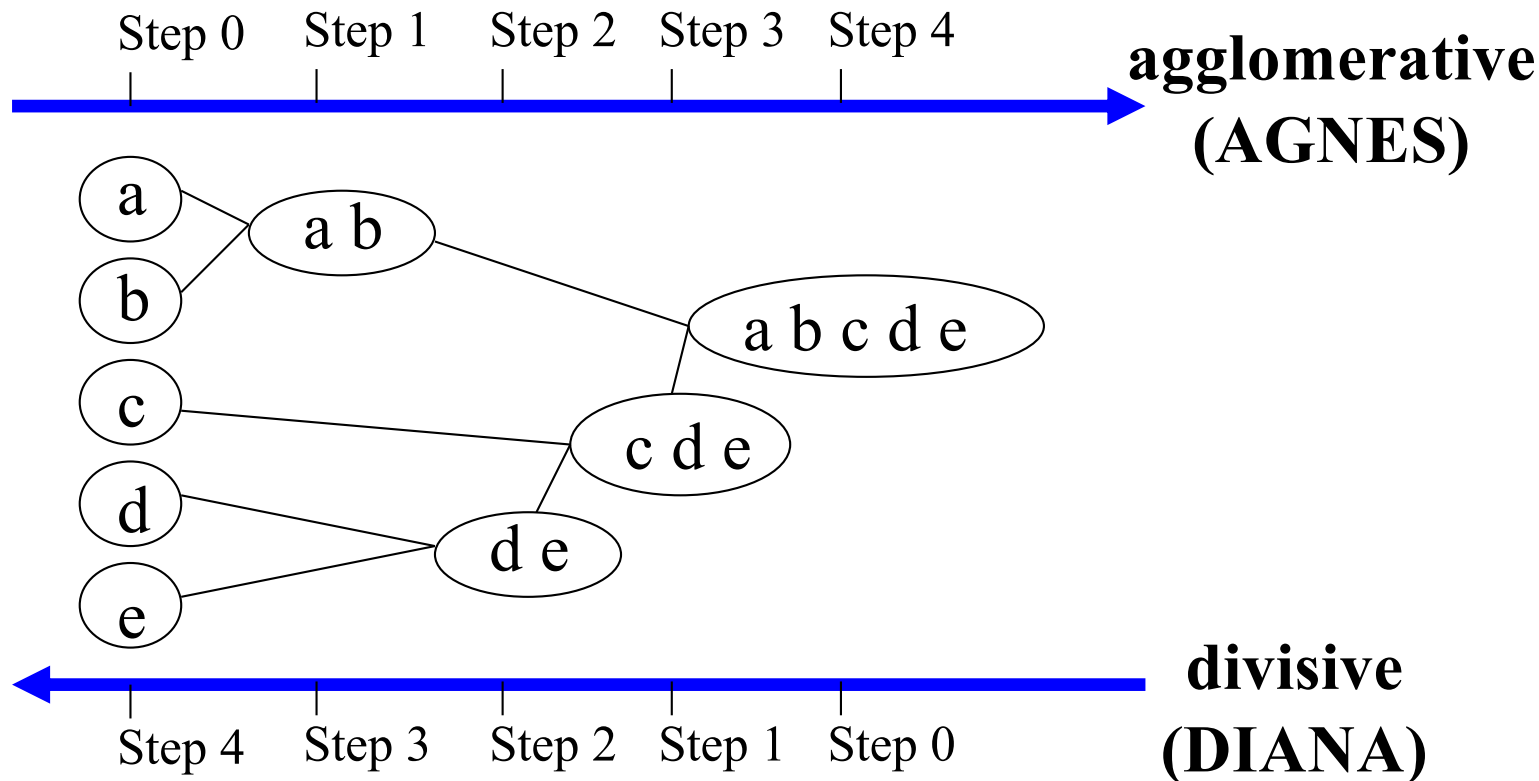
Chapter 7. Cluster Analysis

1. What is Cluster Analysis?
2. Types of Data in Cluster Analysis
3. A Categorization of Major Clustering Methods
4. Partitioning Methods
5. Hierarchical Methods
6. Density-Based Methods
7. Outlier Analysis
8. Summary



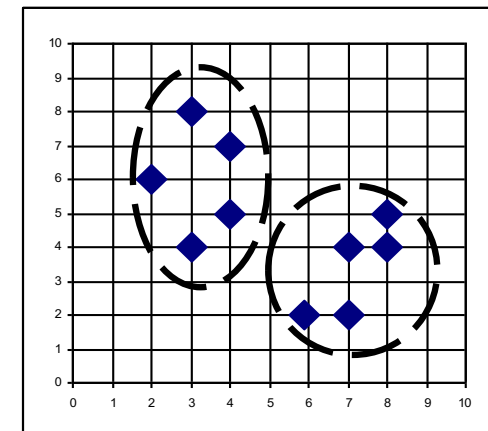
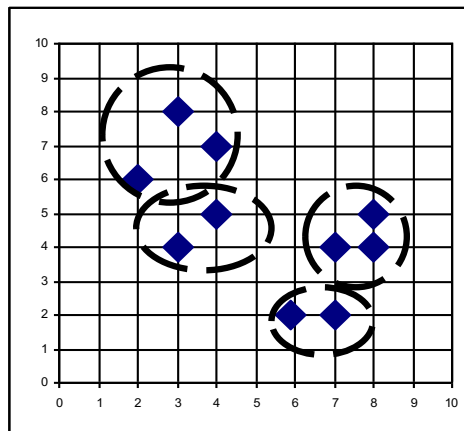
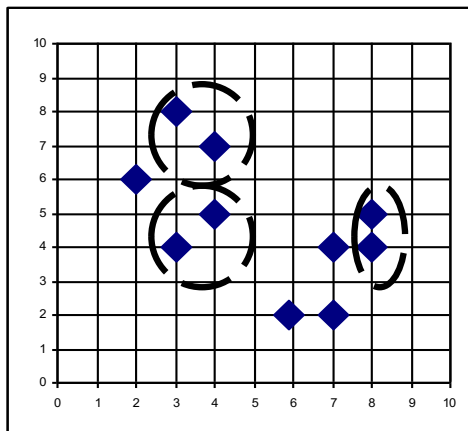
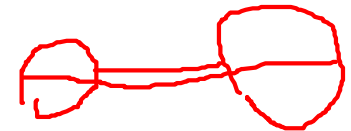
Hierarchical Clustering

- Use a distance matrix as clustering criteria
- Does not require the number of clusters k as an input, but needs a termination condition



AGNES (Agglomerative Nesting)

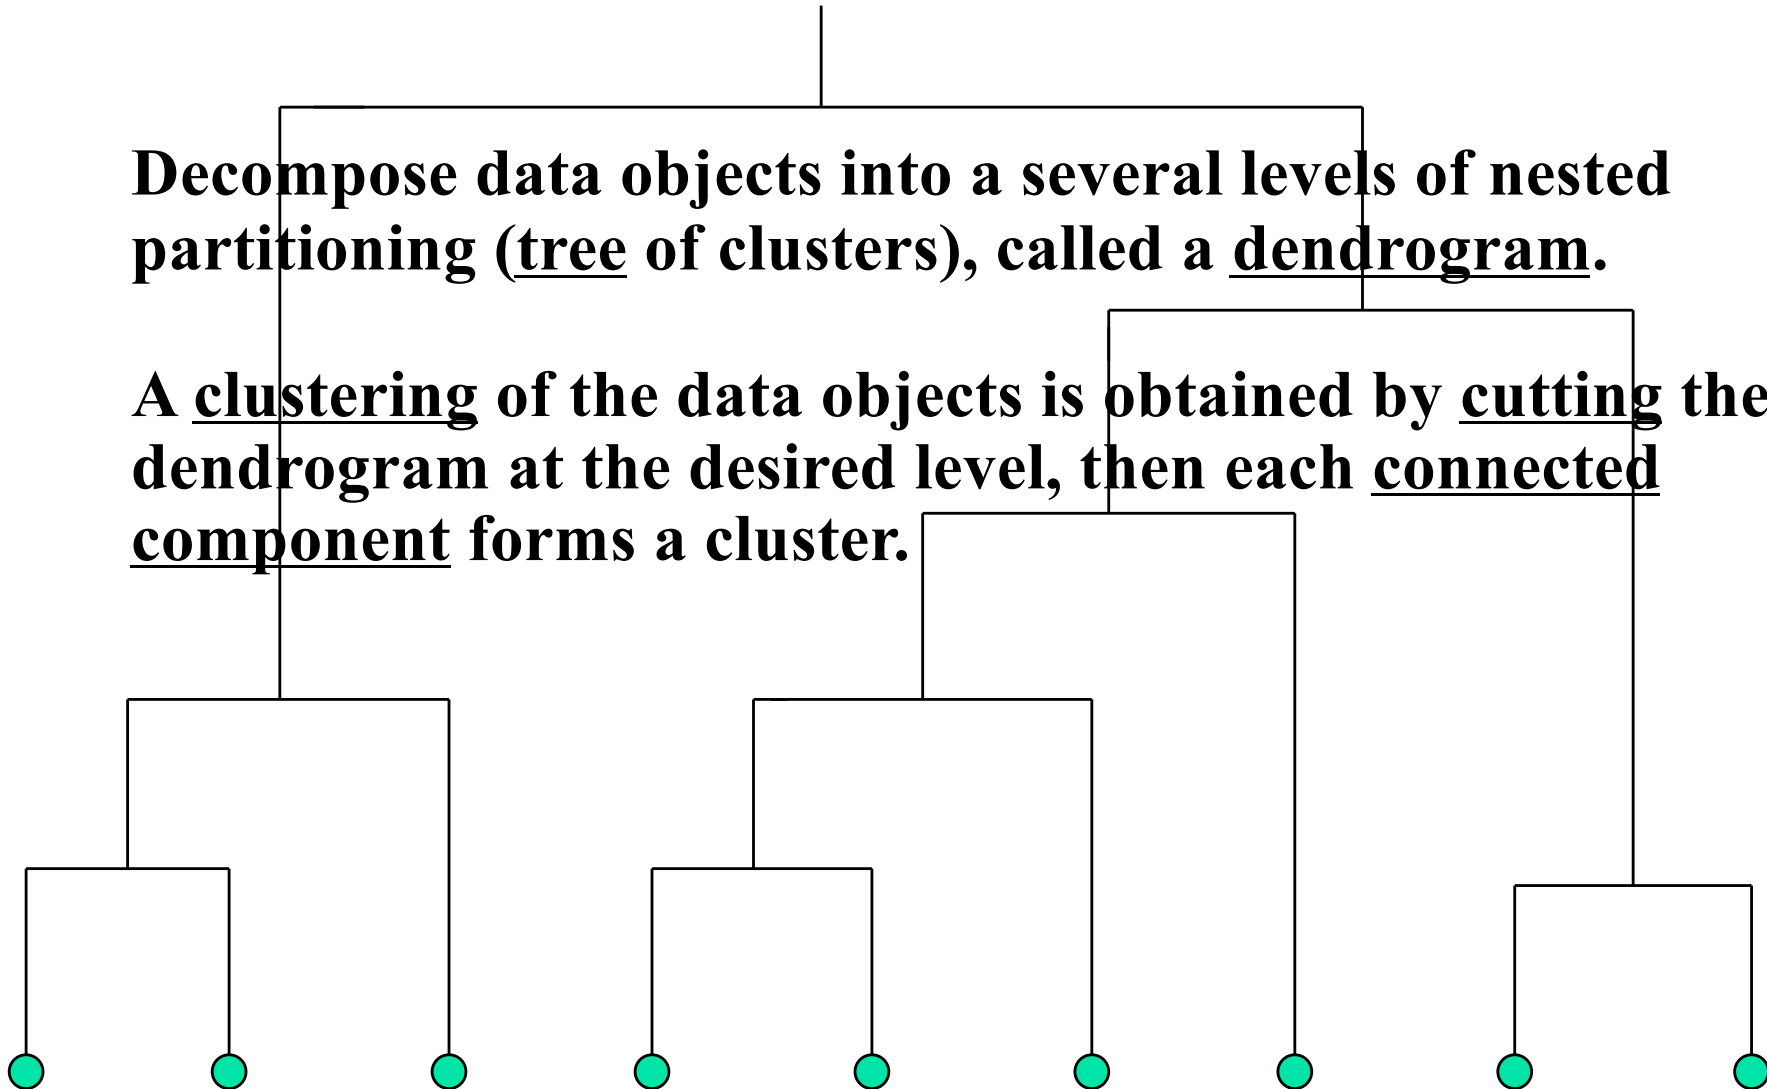
- Introduced in Kaufmann and Rousseeuw (1990)
 - Implemented in statistical analysis packages, Splus
- Use the **single-link** method and the dissimilarity matrix
- Merge nodes that have the least dissimilarity
- Go on in a non-descending fashion
- Eventually all nodes belong to the same cluster



Dendrogram: How the Clusters are Merged

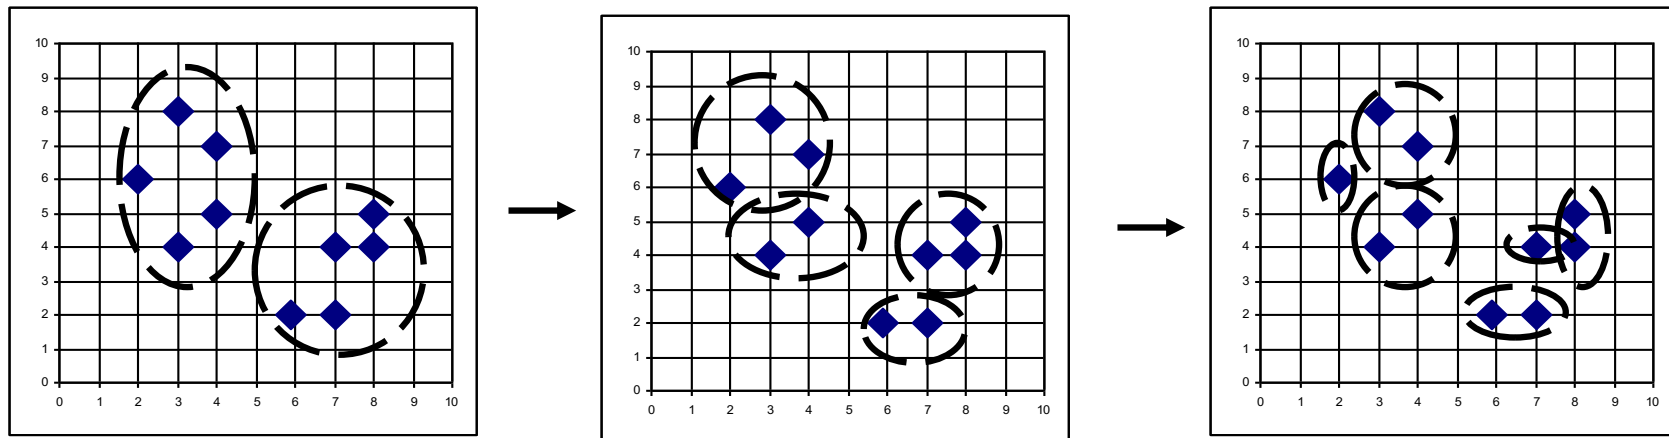
Decompose data objects into a several levels of nested partitioning (tree of clusters), called a dendrogram.

A clustering of the data objects is obtained by cutting the dendrogram at the desired level, then each connected component forms a cluster.



DIANA (Divisive Analysis)

- Introduced in Kaufmann and Rousseeuw (1990)
- Implemented in statistical analysis packages, Splus
- Inverse order of AGNES
- Eventually each node forms a cluster on its own



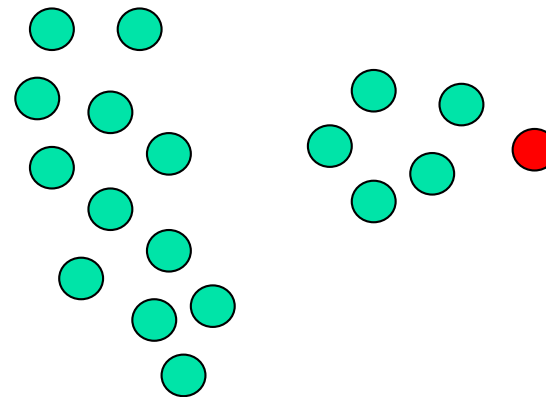
DIANA (Divisive Analysis)

- Outline
 - Initially, there is one large cluster consisting of all n objects
 - At each subsequent step, the largest available cluster is split into two clusters
 - Until finally all clusters comprise of a single object.
 - Thus, the hierarchy is built in $n-1$ steps.
- Complexity in the first step
 - Agglomerative method: $\frac{n(n-1)}{2}$ possible combinations
 - Divisive method: $2^{n-1} - 1$ possible combinations
 - Considerably larger than an agglomerative method



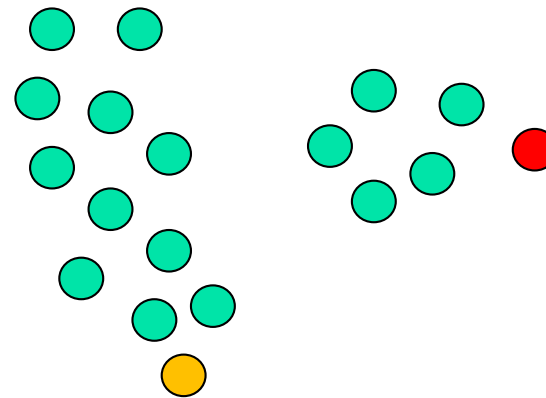
DIANA (Divisive Analysis)

- To avoid considering all possibilities, the algorithm proceeds as follows.
 1. Find the object, which has the highest average dissimilarity to all other objects. This object initiates a new cluster— a sort of a *splinter group*.
 2. For each object i outside the *splinter group*, compute $D_i = [\text{average } d(i, j) \mid j \notin R_{\text{splinter group}}] - [\text{average } d(i, j) \mid j \in R_{\text{splinter group}}]$
 3. Find an object h for which the difference D_h is the largest. If D_h is positive, then h is, on the average close to the splinter group. Put h into the splinter group.



DIANA (Divisive Analysis)

- To avoid considering all possibilities, the algorithm proceeds as follows.
 1. Repeat *Steps* 2 and 3 until all differences D_h are negative. The data set is then split into two clusters.
 2. Select the cluster with the largest **diameter**. The diameter of a cluster is the largest dissimilarity between any two of its objects. Then divide this cluster, following steps 1-4.
 3. Repeat *Step* 5 until all clusters contain only a single object.



Advanced Hierarchical Clustering Methods

- Major weakness of agglomerative clustering methods
 - do not scale well: time complexity of at least $O(n^2)$, where n is the number of total objects
- Integration of hierarchical with distance-based clustering
 - BIRCH (1996): uses CF-tree and incrementally adjusts the quality of sub-clusters
 - ROCK (1999): clustering categorical data by neighbor and link analysis
 - CHAMELEON (1999): hierarchical clustering using dynamic modeling



BIRCH (1996)

- Birch: Balanced Iterative Reducing and Clustering using Hierarchies (Zhang, Ramakrishnan & Livny, SIGMOD'96)
- Incrementally construct a **CF (Clustering Feature)** tree (cf. B-tree), a hierarchical data structure for multiphase clustering
 - Phase 1: scan DB to build an initial in-memory CF tree (a multi-level compression of the data that tries to preserve the inherent clustering structure of the data)
 - Phase 2: use an arbitrary clustering algorithm to cluster the leaf nodes of the CF-tree
- *Scales linearly*: finds a good clustering with a single scan and improves the quality with a few additional scans
- *Weakness*: handles only numeric data, and sensitive to the order of the data records



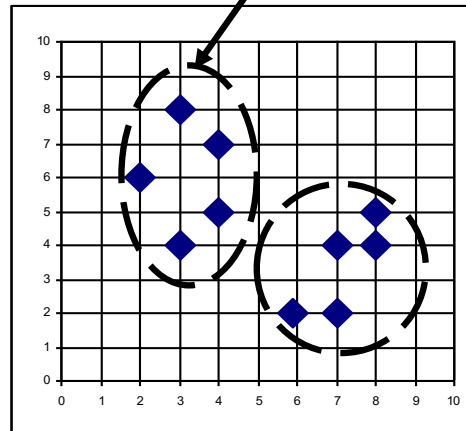
Clustering Feature Vector in BIRCH

Clustering Feature: $CF = (N, \vec{LS}, SS)$

N : **Number of data points**

$$LS: \sum_{i=1}^N \vec{X}_i$$

$$SS: \sum_{i=1}^N \vec{X}_i^2$$



$$CF = (5, (16,30), (54,190))$$

(3,4)

(2,6)

(4,5)

(4,7)

(3,8)



CF-Tree in BIRCH

- Clustering feature:
 - Summary of the statistics for a given cluster: the 0-th, 1st and 2nd moments of the cluster from the statistical point of view
 - Registers crucial measurements for computing cluster and utilizes storage efficiently
- A CF tree is a height-balanced tree that stores the clustering features for a hierarchical clustering
 - A non-leaf node in a tree has descendants or “children”
 - A non-leaf node stores the **sum of the CFs of their children**
- A CF tree has two parameters
 - Branching factor: specify the maximum number of children
 - threshold: max diameter of a cluster stored at the leaf node



The CF Tree Structure

Root

$B = 7$

$L = 6$

CF_1	CF_2	CF_3	CF_6
child ₁	child ₂	child ₃		child ₆

Non-leaf node

CF_1	CF_2	CF_3	CF_5
child ₁	child ₂	child ₃		child ₅

Leaf node

Leaf node

prev	CF_1	CF_2	CF_6	next
------	--------	--------	-------	--------	------

prev	CF_1	CF_2	CF_4	next
------	--------	--------	-------	--------	------

